

# Machine Learning Computer Lab 2 (Group A7)

Qinyuan Qi(qinqi464)      Satya Sai Naga Jaya Koushik Pilla (satpi345)  
Daniele Bozzoli(danbo826)

2024-01-10

## Statement of Contribution

For Lab2, we decided to split the three assignments equally, Qinyuan completed Assignment 1, Satya completed Assignment 2 and Daniele completed Assignment 3, after which, for verification sake, we completed each other's assignments as well and validated our findings. The report was also compiled by three of us, with each handling their respective assignments.

## Assignment 1: Explicit regularization

**Answer:**

(1)

According to the question, we can model create a linear model like the following.

$$Fat = \beta_0 + \beta_1 Channel_1 + \beta_2 Channel_2 + \dots + \beta_{100} Channel_{100} + \epsilon = \sum \beta_i Channel_i + \epsilon$$

In the above formula,  $\beta_0$  is the intercept, while the remaining  $\beta$  are the coefficients corresponding to each channel.

According to the output, the model generated use 100 channel features, and almost all the channels provide contributions to the target, however, p value shows that only very limited channels are useful. And the  $MSE_{test} = 722.4294$ ,  $MSE_{training} = 0.005709117$ . It means training data fit pretty well, however the test data fit not as expected, and the model overfit the data.

```
## test_mse is: 722.4294
## train_mse is: 0.005709117
##
## Call:
## lm(formula = Fat ~ ., data = train_data_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.201500 -0.041315 -0.001041  0.037636  0.187860
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.815e+01  5.488e+00  -3.306  0.01628 *
## Channel1     2.653e+04  1.126e+04   2.357  0.05649 .
## Channel2    -5.871e+04  3.493e+04  -1.681  0.14385
## Channel3     1.154e+05  7.373e+04   1.565  0.16852
## Channel4    -2.432e+05  1.175e+05  -2.070  0.08387 .
```

## Channel5	3.026e+05	1.193e+05	2.536	0.04430	*
## Channel6	-2.365e+05	8.160e+04	-2.898	0.02741	*
## Channel7	1.090e+05	3.169e+04	3.440	0.01380	*
## Channel8	-6.054e+04	1.508e+04	-4.015	0.00700	**
## Channel9	7.871e+04	2.160e+04	3.643	0.01079	*
## Channel10	-1.730e+04	1.640e+04	-1.055	0.33215	
## Channel11	9.562e+04	3.529e+04	2.710	0.03512	*
## Channel12	-2.114e+05	6.198e+04	-3.410	0.01431	*
## Channel13	9.725e+04	4.424e+04	2.198	0.07026	.
## Channel14	5.296e+04	4.666e+04	1.135	0.29968	
## Channel15	-7.855e+04	5.245e+04	-1.498	0.18491	
## Channel16	-8.209e+03	1.893e+04	-0.434	0.67969	
## Channel17	3.769e+04	1.987e+04	1.897	0.10666	
## Channel18	3.306e+04	7.934e+03	4.167	0.00590	**
## Channel19	-8.405e+04	1.929e+04	-4.358	0.00478	**
## Channel20	1.510e+05	3.361e+04	4.492	0.00414	**
## Channel21	-2.069e+05	4.256e+04	-4.862	0.00282	**
## Channel22	1.348e+05	3.824e+04	3.526	0.01243	*
## Channel23	-4.094e+04	3.546e+04	-1.154	0.29222	
## Channel24	2.023e+04	2.761e+04	0.733	0.49134	
## Channel25	3.269e+03	1.071e+04	0.305	0.77045	
## Channel26	-1.297e+04	7.636e+03	-1.699	0.14028	
## Channel27	4.131e+03	1.422e+04	0.291	0.78120	
## Channel28	-4.548e+03	2.988e+04	-0.152	0.88402	
## Channel29	1.089e+04	1.768e+04	0.616	0.56072	
## Channel30	-7.985e+04	2.653e+04	-3.010	0.02371	*
## Channel31	1.756e+05	5.279e+04	3.326	0.01589	*
## Channel32	-1.107e+05	2.904e+04	-3.813	0.00883	**
## Channel33	-6.525e+04	5.407e+04	-1.207	0.27294	
## Channel34	1.007e+05	6.589e+04	1.528	0.17738	
## Channel35	-2.841e+03	1.214e+04	-0.234	0.82266	
## Channel36	-2.268e+04	2.295e+04	-0.988	0.36127	
## Channel37	-4.479e+04	1.292e+04	-3.468	0.01334	*
## Channel38	3.209e+04	1.843e+04	1.742	0.13221	
## Channel39	1.992e+04	2.067e+04	0.964	0.37246	
## Channel40	-9.833e+03	2.431e+04	-0.404	0.69988	
## Channel41	1.659e+04	3.648e+04	0.455	0.66531	
## Channel42	-1.829e+04	3.528e+04	-0.519	0.62260	
## Channel43	-2.423e+04	2.427e+04	-0.998	0.35669	
## Channel44	3.246e+04	2.013e+04	1.613	0.15793	
## Channel45	-8.089e+03	4.023e+04	-0.201	0.84728	
## Channel46	7.065e+03	2.810e+04	0.251	0.80990	
## Channel47	-4.062e+04	1.007e+04	-4.034	0.00685	**
## Channel48	9.080e+04	2.618e+04	3.469	0.01332	*
## Channel49	-6.647e+04	2.372e+04	-2.803	0.03105	*
## Channel50	-4.196e+04	2.856e+04	-1.469	0.19213	
## Channel51	1.097e+05	5.572e+04	1.968	0.09661	.
## Channel52	-1.148e+05	6.376e+04	-1.800	0.12196	
## Channel53	9.525e+04	7.450e+04	1.278	0.24830	
## Channel54	-4.534e+04	7.363e+04	-0.616	0.56067	
## Channel55	-1.535e+03	4.933e+04	-0.031	0.97618	
## Channel56	-2.377e+03	2.109e+04	-0.113	0.91394	
## Channel57	3.174e+04	1.005e+04	3.158	0.01961	*
## Channel58	2.221e+03	1.048e+04	0.212	0.83915	

```

## Channel59 -8.504e+04 2.574e+04 -3.304 0.01634 *
## Channel60 6.382e+04 1.607e+04 3.972 0.00735 **
## Channel61 2.151e+04 1.234e+04 1.742 0.13211
## Channel62 -2.859e+04 1.065e+04 -2.685 0.03631 *
## Channel63 1.796e+04 9.187e+03 1.955 0.09838 .
## Channel64 5.759e+04 3.526e+04 1.633 0.15354
## Channel65 -1.470e+05 6.911e+04 -2.127 0.07752 .
## Channel66 9.121e+04 4.461e+04 2.045 0.08688 .
## Channel67 -5.733e+03 2.197e+04 -0.261 0.80288
## Channel68 -6.290e+04 2.192e+04 -2.870 0.02843 *
## Channel69 6.421e+04 2.074e+04 3.096 0.02121 *
## Channel70 -1.749e+04 1.581e+04 -1.106 0.31111
## Channel71 -7.248e+03 1.934e+04 -0.375 0.72075
## Channel72 3.406e+04 1.185e+04 2.873 0.02830 *
## Channel73 -2.100e+04 1.132e+04 -1.855 0.11308
## Channel74 -3.314e+04 1.220e+04 -2.717 0.03480 *
## Channel75 7.039e+04 2.054e+04 3.427 0.01402 *
## Channel76 -3.187e+04 1.736e+04 -1.836 0.11597
## Channel77 2.061e+04 1.810e+04 1.138 0.29832
## Channel78 -1.180e+04 2.273e+04 -0.519 0.62225
## Channel79 2.669e+04 2.997e+04 0.890 0.40750
## Channel80 -6.051e+04 1.483e+04 -4.080 0.00650 **
## Channel81 1.386e+03 2.628e+04 0.053 0.95966
## Channel82 1.020e+05 4.694e+04 2.173 0.07275 .
## Channel83 -1.706e+05 4.688e+04 -3.640 0.01083 *
## Channel84 1.097e+05 2.892e+04 3.792 0.00905 **
## Channel85 -1.294e+05 3.600e+04 -3.594 0.01145 *
## Channel86 2.130e+05 4.345e+04 4.903 0.00270 **
## Channel87 -1.198e+05 3.818e+04 -3.139 0.02011 *
## Channel88 -2.199e+04 6.085e+04 -0.361 0.73021
## Channel89 7.974e+04 5.077e+04 1.571 0.16733
## Channel90 -1.711e+05 5.499e+04 -3.112 0.02079 *
## Channel91 2.107e+05 6.406e+04 3.289 0.01663 *
## Channel92 -1.959e+05 7.171e+04 -2.733 0.03407 *
## Channel93 2.874e+05 9.937e+04 2.892 0.02762 *
## Channel94 -3.064e+05 9.601e+04 -3.191 0.01881 *
## Channel95 2.048e+05 6.220e+04 3.292 0.01656 *
## Channel96 -5.600e+04 2.929e+04 -1.912 0.10441
## Channel97 -1.318e+04 3.050e+04 -0.432 0.68065
## Channel98 -2.724e+04 2.107e+04 -1.292 0.24375
## Channel99 3.556e+04 1.382e+04 2.573 0.04218 *
## Channel100 -1.206e+04 4.264e+03 -2.828 0.03006 *
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3191 on 6 degrees of freedom
## Multiple R-squared: 1, Adjusted R-squared: 0.9994
## F-statistic: 1651 on 100 and 6 DF, p-value: 1.058e-09

```

(2)

The cost function in lasso regression is:


$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \frac{1}{n} \|X\theta\|^2 + \lambda \|\theta\|_1$$

(3)

According to the plot of `lasso_reg`, we find that when  $\log(\lambda)$  is small, degree of freedom is big which means more features explain the target variable, but when  $\log(\lambda)$  increase, the degree of freedom decrease at the same time, and coefficient get close to 0 in most of the time. But we also find that some of the features's coefficient line fluctuate when it is less than 0. But all of the features's coefficient will converge to 0 at last.

According to the calculation and output of the code, we find that lambda is: 0.05123599  $\log(\lambda)$  is -2.971313 when feature number is 4 (including the intercept).

```
## lambda is 0.05123599 log(lambda) is -2.971313
## when feature number is 4 (including the intercept)
```




(4)

the cost function in ridge regression is:

$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \frac{1}{n} \|X\theta\|^2 + \lambda \|\theta\|_2^2$$

For ridge regression, we find that the degree of freedom always keep same when  $\log(\lambda)$  increase, but the coefficient of features seems to will converge to 0, but after check beta value of `ridge_reg`, we found that the coefficient of features will not converge to 0, but will become a very small value.

Compare to lasso regression, we can not find the required lambda value when degree of freedom is 4, since coefficient value not converge to 0, but converge to a very small value.



(5)

The plot as follows, when  $\log(\lambda)$  increase, CV score increase. Optimized lambda is -5.39019, and in this case, around 10 variables were chosen.

When  $\lambda = -4$ , MSE is higher than when  $\lambda = -5.39019$ . we can say that it results in a statistically significantly better prediction.

The scatter plot as follow.

Since almost all points are around the 45-degree line, it indicates accurate predictions.


```
# when alpha = 1, it is lasso regression
cv_lasso_model <- cv.glmnet(x, y_train, alpha = 1)
summary(cv_lasso_model)
```

```
##          Length Class  Mode
## lambda      77    -none- numeric
## cvm         77    -none- numeric
## cvsd        77    -none- numeric
## cvup        77    -none- numeric
## cvlo        77    -none- numeric
## nzero       77    -none- numeric
## call         4    -none- call
## name         1    -none- character
## glmnet.fit  12    elnet  list
```

```
## lambda.min 1      -none- numeric
## lambda.1se 1      -none- numeric
## index      2      -none- numeric

# Extract lambda values and CV scores
lambda_values <- log(cv_lasso_model$lambda)
cv_scores <- cv_lasso_model$cvm

# Plot the dependence of CV score on log(lambda)
plot(cv_lasso_model)
```




```
plot(lambda_values, cv_scores, type = "b", xlab = "log(lambda)", ylab = "CV Score", main = "CV Score vs lambda")

min_cv_lambda <- log(cv_lasso_model$lambda.min)
min_cv_score <- min(cv_scores)
cat("min_cv_lambda is:", min_cv_lambda, "\n")

## min_cv_lambda is: -5.39019

points(min_cv_lambda, min_cv_score, col = "red", pch = 16, cex = 1.5)
text(min_cv_lambda, min_cv_score, "Minimum CV Score", pos = 3, col = "red")
```


## CV Score vs log(lambda)



```
predictions <- predict(cv_lasso_model, newx = as.matrix(x_test), s = "lambda.min")

plot(y_test, predictions, pch = 16, col = "red",
     xlab = "Original Values", ylab = "Predicted Values",
     main = "Scatter Plot of Test Values for LASSO Model")
abline(a = 0, b = 1, col = "blue", lty = 2)
```

## Scatter Plot of Test Values for LASSO Model



## Assignment 2: Decision trees and logistic regression for bank marketing

Answer:

(1)

```
##### Assignment 2.1 #####
# read data
data <- read.csv("bank-full.csv",header = TRUE, sep = ";")

row_num <- nrow(data)
cols_num <- ncol(data)

# set data split ratio to 0.4, 0.3, 0.3
ratio <- c(train = .4, validate = 0.3, test = .3)

# set random seed
set.seed(12345)

# split data to training,validate and test dataset
train_id <- sample(1:row_num, floor(row_num * ratio[1]))
train_set <- data[train_id, ]

# set random seed
set.seed(12345)
```



```
validation_test_id <- setdiff(1:row_num, train_id)
validation_id <- sample(validation_test_id, floor(row_num * ratio[2]))
validation_set <- data[validation_id, ]

test_id <- setdiff(validation_test_id, validation_id)
test_set <- data[test_id, ]
```

(2)

Decision trees are as follows.

misclassification is listed below.

	misclassification rate of train	misclassification rate of validation
a	0.0995355	0.0995355
b	0.1142446	0.1207697
c	0.06287326	0.1041805

According to the data in the table above, the first one has a good misclassification rate of train and validation. the last one has a good misclassification rate of train but a little bit large misclassification rate on validation set. b will not be considered, since both of values are biggest among 3 methods.

```
##### Assignment 2.2 #####
tree_a <- rpart(y ~ ., data = train_set, method = "class")
tree_b <- rpart(y ~ ., data = train_set, method = "class", control = rpart.control(minbucket = 7000))
tree_c <- rpart(y ~ ., data = train_set, method = "class", control = rpart.control(cp = 0.0005))

train_predictions_a <- predict(tree_a, train_set, type = "class")
validation_predictions_a <- predict(tree_a, validation_set, type = "class")

train_predictions_b <- predict(tree_b, train_set, type = "class")
validation_predictions_b <- predict(tree_b, validation_set, type = "class")

train_predictions_c <- predict(tree_c, train_set, type = "class")
validation_predictions_c <- predict(tree_c, validation_set, type = "class")

# misclassification rates
train_misclassification_rate_a <- mean(train_predictions_a != train_set$y)
validation_misclassification_rate_a <- mean(validation_predictions_a != validation_set$y)

train_misclassification_rate_b <- mean(train_predictions_b != train_set$y)
validation_misclassification_rate_b <- mean(validation_predictions_b != validation_set$y)

train_misclassification_rate_c <- mean(train_predictions_c != train_set$y)
validation_misclassification_rate_c <- mean(validation_predictions_c != validation_set$y)

cat("train_misclassification_rate_a is ",train_misclassification_rate_a,"\n")

## train_misclassification_rate_a is 0.0995355
cat("validation_misclassification_rate_a is ",validation_misclassification_rate_a,"\n")

## validation_misclassification_rate_a is 0.0995355
```

```

cat("train_misclassification_rate_b is ",train_misclassification_rate_b,"\n")

## train_misclassification_rate_b is  0.1142446

cat("validation_misclassification_rate_b is ",validation_misclassification_rate_b,"\n")

## validation_misclassification_rate_b is  0.1207697

cat("train_misclassification_rate_c is ",train_misclassification_rate_c,"\n")

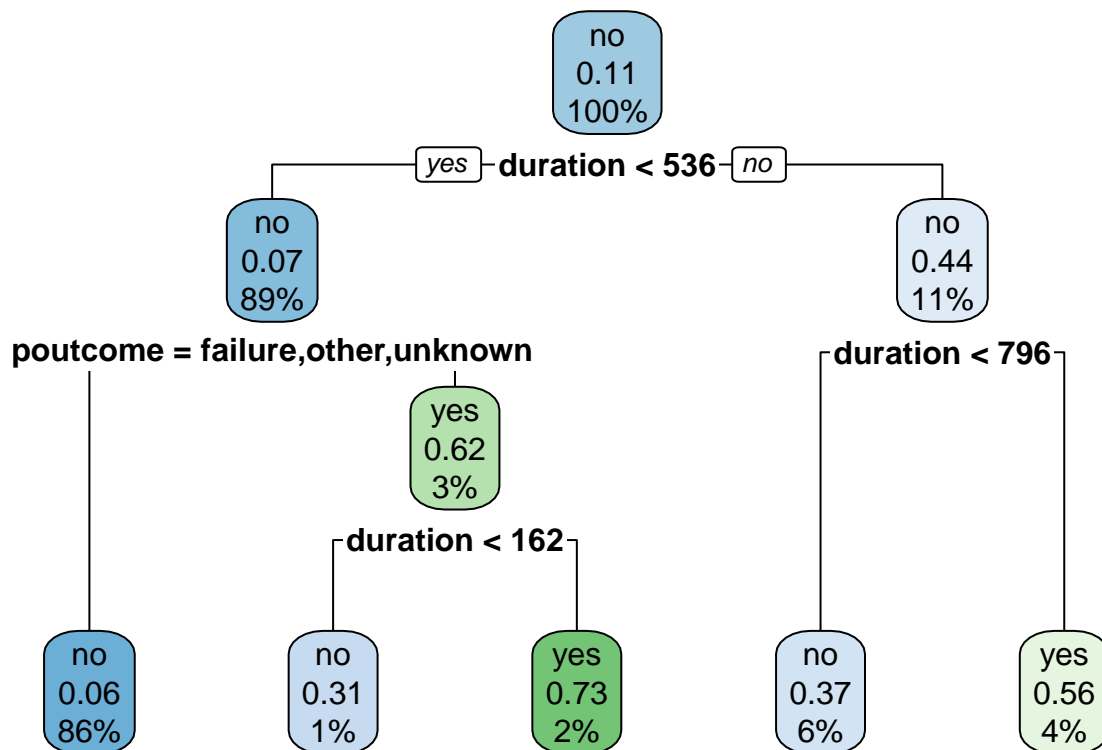
## train_misclassification_rate_c is  0.06287326

cat("validation_misclassification_rate_c is ",validation_misclassification_rate_c,"\n")

## validation_misclassification_rate_c is  0.1041805

rpart.plot(tree_a)

```



```


rpart.plot(tree_b)

```

no  
0.11  
100%

```
rpart.plot(tree_c)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



(3)

```
##### Assignment 2.3 #####
# Set folds number
num_folds <- 5
cv_inds <- cut(seq(1, nrow(train_set)), breaks = num_folds, labels = FALSE)
cv_errors <- numeric(30)
train_deviances_depth <- numeric(30)
validation_deviances_depth <- numeric(30)

for (depth in 1:30) {
  tree <- rpart(y ~ ., data = train_set, method = "class",
    control = rpart.control(cp = 0.0005), maxdepth = depth)

  validation_predictions <- predict(tree, validation_set, type = "class")

  cv_errors_depth <- mean(validation_predictions != validation_set$y)

  cv_errors[depth] <- cv_errors_depth

  train_deviances_depth[depth] <- sum(deviance(tree))
  validation_deviances_depth[depth] <- sum(deviance(tree, newdata = validation_set))
}

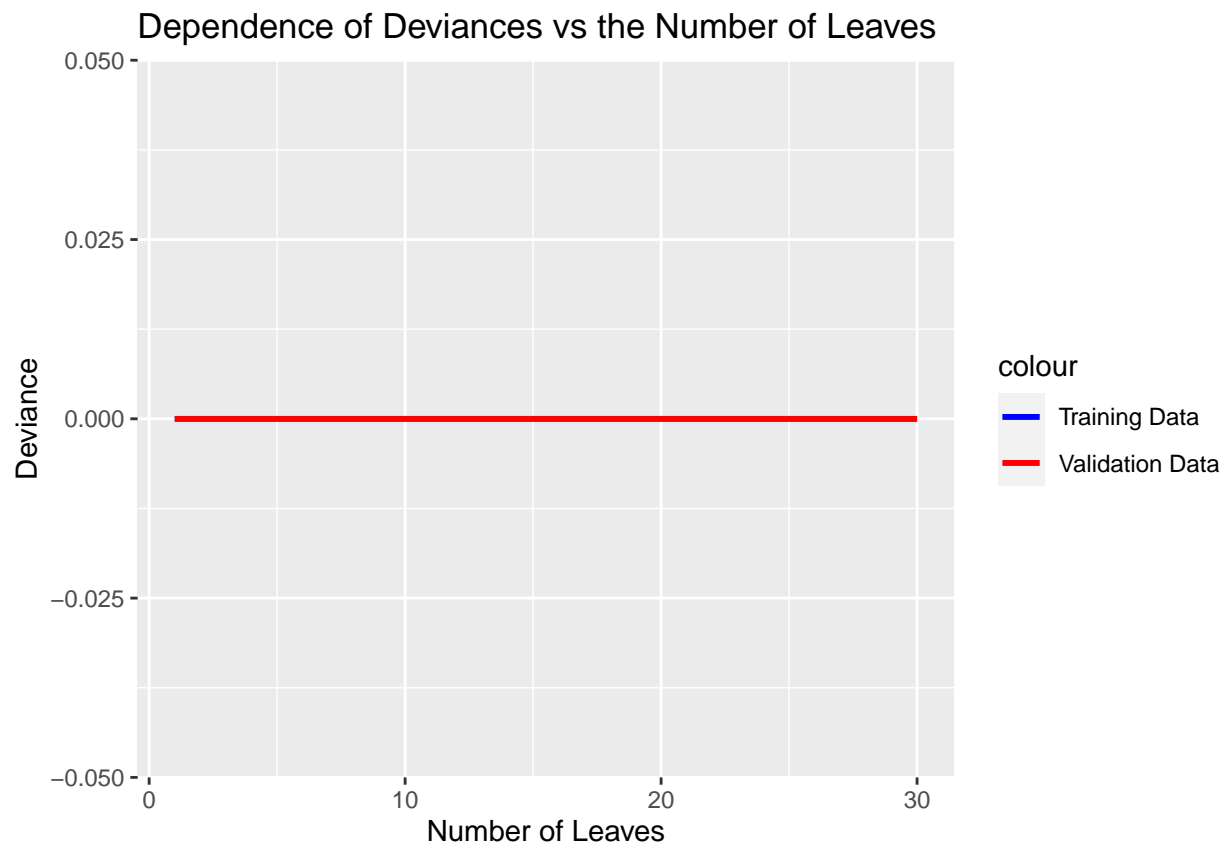
optimal_depth <- which.min(cv_errors)
```

```
cat("Optimal tree depth:", optimal_depth, "\n")

## Optimal tree depth: 1

plot_data <- data.frame(Leaves = 1:30, Train_Deviance = train_deviances_depth,
  Validation_Deviance = validation_deviances_depth)

ggplot(plot_data, aes(x = Leaves)) +
  geom_line(aes(y = Train_Deviance, color = "Training Data"), size = 1) +
  geom_line(aes(y = Validation_Deviance, color = "Validation Data"), size = 1) +
  labs(title = "Dependence of Deviances vs the Number of Leaves",
    x = "Number of Leaves",
    y = "Deviance") +
  scale_color_manual(values = c("Training Data" = "blue", "Validation Data" = "red"))
```



(4)

An optimized tree depth = 1.

```
##### Assignment 2.4 #####
final_tree <- rpart(y ~ ., data = train_set, method = "class",
  control = rpart.control(cp = 0.0005), maxdepth = 1)

validation_predictions <- predict(final_tree, validation_set, type = "class")

conf_matrix <- confusionMatrix(validation_predictions, validation_set$y)
```

```

accuracy <- conf_matrix$overall["Accuracy"]
f1_score <- conf_matrix$byClass["F1"]

# Print the confusion matrix, accuracy, and F1 score
cat("Confusion Matrix:\n")
cat(conf_matrix$table, "\n")

cat("Accuracy:", accuracy, "\n")
cat("F1 Score:", f1_score, "\n")

```

(5)

```

##### Assignment 2.5 #####
loss_matrix <- matrix(c(0, 1, 5, 0), nrow = 2)
test_predictions <- predict(final_tree, test_set, type = "class", parms = list(loss = loss_matrix))
conf_matrix <- table(Actual = test_set$y, Predicted = test_predictions)
cat("Confusion Matrix:\n")

```

## Confusion Matrix:

```
print(conf_matrix)
```

```

##      Predicted
## Actual    no   yes
##    no 11471  508
##    yes  928  657

```

(6)

```

##### Assignment 2.6 #####
threshold <- 0.05
logistic_model <- glm(y ~ ., data = train_set, family = "binomial")

logistic_probabilities <- predict(logistic_model, test_set, type = "response")
logistic_predictions <- ifelse(logistic_probabilities > threshold, "yes", "no")

conf_matrix_logistic <- confusionMatrix(logistic_predictions, test_set$y)

TP <- conf_matrix_logistic$table["yes", "yes"]
FN <- conf_matrix_logistic$table["no", "yes"]
FP <- conf_matrix_logistic$table["yes", "no"]
TN <- conf_matrix_logistic$table["no", "no"]

TPR <- TP / (TP + FN)
FPR <- FP / (FP + TN)

# Print the results
cat("True Positive Rate (TPR):", TPR, "\n")
cat("False Positive Rate (FPR):", FPR, "\n")

```

### Assignment 3. Principal components and implicit regularization

Answer:

(1)

According to the output, we need 35 components to obtain at least 95% of variance in the data. The proportion of variation explained by first and second principal components are 0.2501699 and 0.1693597 respectively.

```
##### Assignment 3.1 #####
n <- nrow(data)
features <- data[, -101]

s_features <- scale(features)

S <- (t(s_features) %*% s_features)/n # sample covariance matrix

Eig <- eigen(S)

# eigen in descending order
s_indx <- order(Eig$values, decreasing = TRUE)
s_eig <- Eig$values[s_indx]

# cumulative explained variance
cum_var <- cumsum(s_eig) / sum(s_eig)

q_95 <- which(cum_var >= 0.95)[1] # q for 95% var

first_two_components <- Eig$vectors[, 1:2] # first two PC

# proportion of variation explained by each of the first two components
PC1_var <- s_eig[1] / sum(s_eig)
PC2_var <- s_eig[2] / sum(s_eig)


cat("Number of components needed for 95% variance:", q_95, "
Proportion of variation explained by the first component:", PC1_var, "
Proportion of variation explained by the second component:", PC2_var, "\n")

## Number of components needed for 95% variance: 35
## Proportion of variation explained by the first component: 0.2501699
## Proportion of variation explained by the second component: 0.1693597
```

(2)


We observe from the plot how the variable “state”, the first feature of PC1, explains most of the data. The other variables carry significantly less explanation. The other 4 features that contribute the most are features 93 (PctBornSameState), 61 (PctSpeakEnglOnly), 5 (racePctWhite) and 76 (PersPerRentOccHous).

```
##### Assignment 3.2 #####
PCA <- princomp(features)
L <- PCA$loadings
plot(L[,1], xlim=c(0,20))
```



```
highest5 <- order(L[,1], decreasing=TRUE)[1:5]  
plot(L[,1], L[,2], main="PCA Scores")
```





(3)

Train MSE is 0.2752071 and Test MES is 0.4248011. We observe how the MSE of the test data is significantly bigger than the MSE obtained from train data. This might happen because the model chosen is overfitting our given data.

```
##### Assignment 3.3 #####
# train and test data
id <- sample(1:n, floor(n*0.5))
trn <- data[id,]
tst <- data[-id,]

# scaling
scaler <- preProcess(trn)
trainS <- predict(scaler,trn)
testS <- predict(scaler,tst)

# linear regression model and test data predictions
linmod <- lm(trainS$ViolentCrimesPerPop ~ ., trainS)
test_pred <- predict(linmod, testS[, -101])

# training and test data MSE
train_MSE <- mean((trainS$ViolentCrimesPerPop - linmod$fitted.values)^2)
test_MSE <- mean((testS$ViolentCrimesPerPop - test_pred)^2)

cat("Train mean squared error:", train_MSE, "\nTest mean squared error:", test_MSE)
```

```
## Train mean squared error: 0.2752071
## Test mean squared error: 0.4248011
```

(4)

Looking at the plots we can see how the errors converge to zero after 1700 iterations (approx 1200 in the second graph but considering that we took off the first 500). Hence 1700 is the optimal iteration number to get good results. The following iterations do not significantly improve our model and hence can lead us to overfitting.

```
##### Assignment 3.4 #####
train_new <- as.matrix(trainS[,-101]) # training data - response variable
train_r <- trainS[,101] # training response variable
```

```
test_new <- as.matrix(testS[,-101]) # test data - response variable
test_r <- trainS[,101] # test response variable
```

```
# error vectors for training and test data
train_e <- c()
test_e <- c()
```

```
costfun <- function(theta_vec){
  train_cost <- mean(((train_new %*% theta_vec) - train_r)^2)
  train_e <- c(train_e, train_cost)
  test_cost <- mean(((test_new %*% theta_vec) - test_r)^2)
  test_e <- c(test_e, test_cost)
  return(train_cost)
}
```

```
theta0 <- rep(0,100)
opt <- optim(theta0, method="BFGS", costfun)

opt$val
```


```
## [1] 0.2752213
```

```
par(mfrow=c(1,2))
```


```
plot(ylab="errors", xlab="iteration", train_e, type = "l", main="Errors")
lines(test_e, col="red")
legend("topright", legend = c("Train errors", "Test errors"),
col = c("black", "red"), lty = 1, cex = 0.5)
```

```
plot(ylab="errors", xlab="iteration", train_e[-c(1:500)][1:2000],
type = "l", main="Zoom on errors")
lines(test_e[-c(1:500)][1:2000], col="red")
legend("topright", legend = c("Train errors", "Test errors"),
col = c("black", "red"), lty = 1, cex = 0.5)
```

### Errors



### Zoom on errors



## Appendix: All code for this report

```
##### Init code For Assignment 1 #####
rm(list = ls())
library(plyr)
library(readr)
library(dplyr)
library(caret)
library(ggplot2)
library(repr)
library(glmnet)
library(rpart)
library(rpart.plot)

# set random seed
set.seed(12345)
##### Assignment 1.1 #####

# read data
data <- read.csv("tecator.csv")
data <- data %>% select(Fat,Channel1:Channel100)

row_num <- nrow(data)
# set data split ratio to 0.5, 0.5
ratio <- c(train = .5, test = .5)

# split data to training and test dataset
train_id <- sample(1:row_num, floor(row_num * ratio[1]))
train_data_set <- data[train_id, ]

test_id <- setdiff(1:row_num, train_id)
test_data_set <- data[test_id, ]

# fit linear model
lm_model <- lm(Fat ~ ., data = train_data_set)

predicted_fat_test <- predict(lm_model, test_data_set)
predicted_fat_train <- predict(lm_model, train_data_set)

# calc the mean value
test_mse <- mean((predicted_fat_test - test_data_set$Fat)^2)
train_mse <- mean((predicted_fat_train - train_data_set$Fat)^2)
cat("test_mse is:",test_mse,"\n")
cat("train_mse is:",train_mse,"\n")

summary(lm_model)
##### Assignment 1.3 #####
# regularize the data before we fit the model
pre_proc_val <- preProcess(train_data_set, method = c("center", "scale"))
train_data_set <- predict(pre_proc_val, train_data_set)
test_data_set <- predict(pre_proc_val, test_data_set)

x_train <- as.matrix(train_data_set %>% select(-Fat))
```

```

y_train <- as.matrix(train_data_set %>% select(Fat))

# when alpha = 1, it is lasso regression
lasso_reg <- glmnet(x_train, y_train, alpha = 1, family = "gaussian")

#summary(lasso_reg)

lambda_deg_freedom <- data.frame(lambda = lasso_reg$lambda,
                                deg_freedom = lasso_reg$df)

lambda_deg_freedom <- lambda_deg_freedom %>% arrange(desc(deg_freedom))

lambda_deg_freedom_num <- lambda_deg_freedom[which(lambda_deg_freedom$deg_freedom == 4),] %>%
  head(1) %>% pull(lambda)

cat("lambda is ", lambda_deg_freedom_num,
    " log(lambda) is ", log(lambda_deg_freedom_num),
    "\nwhen feature number is 4 (including the intercept)\n")

plot(lasso_reg, xvar="lambda", xlab='Log Lambda')
abline(v=log(lambda_deg_freedom_num), col="red")
##### Assignment 1.4 #####
# when alpha = 0, it is ridge regression
ridge_reg <- glmnet(x_train, y_train, alpha = 0, family = "gaussian")

lambda_deg_freedom <- data.frame(lambda = ridge_reg$lambda,
                                deg_freedom = ridge_reg$df)

lambda_deg_freedom <- lambda_deg_freedom %>% arrange(desc(deg_freedom))

lambda_deg_freedom_num <- lambda_deg_freedom[which(lambda_deg_freedom$deg_freedom == 4),] %>%
  head(1) %>% pull(lambda)

plot(ridge_reg, xvar="lambda", xlab='Log Lambda')
# when alpha = 1, it is lasso regression
cv_lasso_model <- cv.glmnet(x, y_train, alpha = 1)
summary(cv_lasso_model)

# Extract lambda values and CV scores
lambda_values <- log(cv_lasso_model$lambda)
cv_scores <- cv_lasso_model$cvm

# Plot the dependence of CV score on log(lambda)
plot(cv_lasso_model)
plot(lambda_values, cv_scores, type = "b", xlab = "log(lambda)", ylab = "CV Score", main = "CV Score vs ")

min_cv_lambda <- log(cv_lasso_model$lambda.min)
min_cv_score <- min(cv_scores)
cat("min_cv_lambda is:", min_cv_lambda, "\n")
points(min_cv_lambda, min_cv_score, col = "red", pch = 16, cex = 1.5)
text(min_cv_lambda, min_cv_score, "Minimum CV Score", pos = 3, col = "red")

```

```

predictions <- predict(cv_lasso_model, newx = as.matrix(x_test), s = "lambda.min")

plot(y_test, predictions, pch = 16, col = "red",
     xlab = "Original Values", ylab = "Predicted Values",
     main = "Scatter Plot of Test Values for LASSO Model")
abline(a = 0, b = 1, col = "blue", lty = 2)
##### Init code For Assignment 1 #####
rm(list = ls())

##### Assignment 2.1 #####
# read data
data <- read.csv("bank-full.csv", header = TRUE, sep = ";")

row_num <- nrow(data)
cols_num <- ncol(data)

# set data split ratio to 0.4, 0.3, 0.3
ratio <- c(train = .4, validate = 0.3, test = .3)

# set random seed
set.seed(12345)

# split data to training, validate and test dataset
train_id <- sample(1:row_num, floor(row_num * ratio[1]))
train_set <- data[train_id, ]

# set random seed
set.seed(12345)

validation_test_id <- setdiff(1:row_num, train_id)
validation_id <- sample(validation_test_id, floor(row_num * ratio[2]))
validation_set <- data[validation_id, ]

test_id <- setdiff(validation_test_id, validation_id)
test_set <- data[test_id, ]

##### Assignment 2.2 #####
tree_a <- rpart(y ~ ., data = train_set, method = "class")
tree_b <- rpart(y ~ ., data = train_set, method = "class", control = rpart.control(minbucket = 7000))
tree_c <- rpart(y ~ ., data = train_set, method = "class", control = rpart.control(cp = 0.0005))

train_predictions_a <- predict(tree_a, train_set, type = "class")
validation_predictions_a <- predict(tree_a, validation_set, type = "class")

train_predictions_b <- predict(tree_b, train_set, type = "class")
validation_predictions_b <- predict(tree_b, validation_set, type = "class")

train_predictions_c <- predict(tree_c, train_set, type = "class")
validation_predictions_c <- predict(tree_c, validation_set, type = "class")

# misclassification rates
train_misclassification_rate_a <- mean(train_predictions_a != train_set$y)
validation_misclassification_rate_a <- mean(validation_predictions_a != validation_set$y)

```

```

train_misclassification_rate_b <- mean(train_predictions_b != train_set$y)
validation_misclassification_rate_b <- mean(validation_predictions_b != validation_set$y)

train_misclassification_rate_c <- mean(train_predictions_c != train_set$y)
validation_misclassification_rate_c <- mean(validation_predictions_c != validation_set$y)

cat("train_misclassification_rate_a is ",train_misclassification_rate_a,"\n")
cat("validation_misclassification_rate_a is ",validation_misclassification_rate_a,"\n")

cat("train_misclassification_rate_b is ",train_misclassification_rate_b,"\n")
cat("validation_misclassification_rate_b is ",validation_misclassification_rate_b,"\n")

cat("train_misclassification_rate_c is ",train_misclassification_rate_c,"\n")
cat("validation_misclassification_rate_c is ",validation_misclassification_rate_c,"\n")

rpart.plot(tree_a)
rpart.plot(tree_b)
rpart.plot(tree_c)
##### Assignment 2.3 #####
# Set folds number
num_folds <- 5
cv_inds <- cut(seq(1, nrow(train_set)), breaks = num_folds, labels = FALSE)
cv_errors <- numeric(30)
train_deviances_depth <- numeric(30)
validation_deviances_depth <- numeric(30)

for (depth in 1:30) {
  tree <- rpart(y ~ ., data = train_set, method = "class",
    control = rpart.control(cp = 0.0005), maxdepth = depth)

  validation_predictions <- predict(tree, validation_set, type = "class")

  cv_errors_depth <- mean(validation_predictions != validation_set$y)

  cv_errors[depth] <- cv_errors_depth

  train_deviances_depth[depth] <- sum(deviance(tree))
  validation_deviances_depth[depth] <- sum(deviance(tree, newdata = validation_set))
}

optimal_depth <- which.min(cv_errors)
cat("Optimal tree depth:", optimal_depth, "\n")

plot_data <- data.frame(Leaves = 1:30, Train_Deviance = train_deviances_depth,
  Validation_Deviance = validation_deviances_depth)

ggplot(plot_data, aes(x = Leaves)) +
  geom_line(aes(y = Train_Deviance, color = "Training Data"), size = 1) +
  geom_line(aes(y = Validation_Deviance, color = "Validation Data"), size = 1) +
  labs(title = "Dependence of Deviances vs the Number of Leaves",
    x = "Number of Leaves",

```

```

    y = "Deviance") +
    scale_color_manual(values = c("Training Data" = "blue", "Validation Data" = "red"))
##### Assignment 2.4 #####
final_tree <- rpart(y ~ ., data = train_set, method = "class",
control = rpart.control(cp = 0.0005), maxdepth = 1)

validation_predictions <- predict(final_tree, validation_set, type = "class")

conf_matrix <- confusionMatrix(validation_predictions, validation_set$y)

accuracy <- conf_matrix$overall["Accuracy"]
f1_score <- conf_matrix$byClass["F1"]

# Print the confusion matrix, accuracy, and F1 score
cat("Confusion Matrix:\n")
cat(conf_matrix$table, "\n")

cat("Accuracy:", accuracy, "\n")
cat("F1 Score:", f1_score, "\n")
##### Assignment 2.5 #####
loss_matrix <- matrix(c(0, 1, 5, 0), nrow = 2)
test_predictions <- predict(final_tree, test_set, type = "class", parms = list(loss = loss_matrix))
conf_matrix <- table(Actual = test_set$y, Predicted = test_predictions)
cat("Confusion Matrix:\n")
print(conf_matrix)

##### Assignment 2.6 #####
threshold <- 0.05
logistic_model <- glm(y ~ ., data = train_set, family = "binomial")

logistic_probabilities <- predict(logistic_model, test_set, type = "response")
logistic_predictions <- ifelse(logistic_probabilities > threshold, "yes", "no")

conf_matrix_logistic <- confusionMatrix(logistic_predictions, test_set$y)

TP <- conf_matrix_logistic$table["yes", "yes"]
FN <- conf_matrix_logistic$table["no", "yes"]
FP <- conf_matrix_logistic$table["yes", "no"]
TN <- conf_matrix_logistic$table["no", "no"]

TPR <- TP / (TP + FN)
FPR <- FP / (FP + TN)

# Print the results
cat("True Positive Rate (TPR):", TPR, "\n")
cat("False Positive Rate (FPR):", FPR, "\n")
##### Init code For Assignment 3 #####
rm(list=ls(all.names = T))
library(ggplot2)
library(caret)
set.seed(12345)
data <- read.csv("communities.csv")
##### Assignment 3.1 #####

```



```

n <- nrow(data)
features <- data[, -101]

s_features <- scale(features)

S <- (t(s_features) %*% s_features)/n # sample covariance matrix

Eig <- eigen(S)

# eigen in descending order
s_indx <- order(Eig$values, decreasing = TRUE)
s_eig <- Eig$values[s_indx]

# cumulative explained variance
cum_var <- cumsum(s_eig) / sum(s_eig)

q_95 <- which(cum_var >= 0.95)[1] # q for 95% var

first_two_components <- Eig$vectors[, 1:2] # first two PC

# proportion of variation explained by each of the first two components
PC1_var <- s_eig[1] / sum(s_eig)
PC2_var <- s_eig[2] / sum(s_eig)

cat("Number of components needed for 95% variance:", q_95, "
Proportion of variation explained by the first component:", PC1_var, "
Proportion of variation explained by the second component:", PC2_var, "\n")
##### Assignment 3.2 #####
PCA <- princomp(features)
L <- PCA$loadings
plot(L[,1], xlim=c(0,20))
highest5 <- order(L[,1], decreasing=TRUE)[1:5]
plot(L[,1], L[,2], main="PCA Scores")
##### Assignment 3.3 #####
# train and test data
id <- sample(1:n, floor(n*0.5))
trn <- data[id,]
tst <- data[-id,]

# scaling
scaler <- preProcess(trn)
trainS <- predict(scaler,trn)
testS <- predict(scaler,tst)

# linear regression model and test data predictions
linmod <- lm(trainS$ViolentCrimesPerPop ~ ., trainS)
test_pred <- predict(linmod, testS[, -101])

# training and test data MSE
train_MSE <- mean((trainS$ViolentCrimesPerPop - linmod$fitted.values)^2)
test_MSE <- mean((testS$ViolentCrimesPerPop - test_pred)^2)

cat("Train mean squared error:", train_MSE, "\nTest mean squared error:", test_MSE)

```

```
##### Assignment 3.4 #####
train_new <- as.matrix(trainS[,-101]) # training data - response variable
train_r <- trainS[,101] # training response variable

test_new <- as.matrix(testS[,-101]) # test data - response variable
test_r <- trainS[,101] # test response variable

# error vectors for training and test data
train_e <- c()
test_e <- c()

costfun <- function(theta_vec){
  train_cost <- mean(((train_new %*% theta_vec) - train_r)^2)
  train_e <- c(train_e, train_cost)
  test_cost <- mean(((test_new %*% theta_vec) - test_r)^2)
  test_e <- c(test_e, test_cost)
  return(train_cost)
}

theta0 <- rep(0,100)
opt <- optim(theta0, method="BFGS", costfun)

opt$val

par(mfrow=c(1,2))

plot(ylab="errors", xlab="iteration",train_e, type = "l", main="Errors")
lines(test_e, col="red")
legend("topright", legend = c("Train errors", "Test errors"),
col = c("black", "red"), lty = 1, cex = 0.5)

plot(ylab="errors", xlab="iteration",train_e[-c(1:500)][1:2000],
type = "l", main="Zoom on errors")
lines(test_e[-c(1:500)][1:2000], col="red")
legend("topright", legend = c("Train errors", "Test errors"),
col = c("black", "red"), lty = 1, cex = 0.5)
```